

Kokkos Kernels Math Library

Luc Berger-Vergiat, S. Rajamanickam, V. Dang,
N. Ellingwood, J. Foucar, E. Harvey, B. Kelley,
K. Liegeois, J. Loe, C. Pearson

ECP Annual Meeting

May 5th 2022

Sandia National Laboratories is a multi-mission laboratory managed and operated by National Technology and Engineering Solutions of Sandia, LLC., a wholly owned subsidiary of Honeywell International, Inc., for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA0003525.
SAND2022-6041 C

The aims of Kokkos Kernels are to:

- ▶ deliver **portable** sparse/dense linear algebra and graph kernels,
- ▶ deliver **robust software ecosystem** for other software technology projects and applications,
- ▶ serve as **reference implementation** of key kernel needs of applications,
- ▶ partner with libraries, applications and vendors to identify new opportunities for performance.

Major partners and customers: Trilinos, PETSc, ExaWind, ORNL, ANL, QMCPACK, Nvidia, Intel, AMD

- ▶ <https://github.com/kokkos/kokkos-kernels>:
 - ▶ Kokkos Kernels GitHub repository,
 - ▶ <https://github.com/kokkos/kokkos-kernels/wiki>,
 - ▶ The wiki provides API calls, examples and build instructions.
- ▶ <https://kokkosteam.slack.com>:
 - ▶ Slack workspace for Kokkos, includes a kokkos-kernels channel,
 - ▶ Please join: fastest way to get your questions answered.

A focus on device BLAS and batched BLAS kernels

Learning objectives:

- ▶ Motivation for batched functions
- ▶ Two namespaces with BLAS and LAPACK functions
- ▶ Calling batched functions

KokkosBlas namespace

- ▶ **KokkosBlas:** device and functor level functions
 - ▶ *Intended Use Case:*
 - ▶ Caller uses optimal amount of parallelism to work on single input data
 - ▶ *Multiple Interfaces: Serial, Team, TeamVector, Device*
 - ▶ Device: all levels of nested parallelism are used on whole device
 - ▶ TeamVector: two-level nested parallelism is used with *TeamThreadRange* and *TeamVectorRange*
 - ▶ Team: one-level nested parallelism is used with *TeamThreadRange*
 - ▶ Serial: no nested parallelism is used internally

KokkosBatched namespace

- ▶ **KokkosBatched:** functor level functions
 - ▶ *Intended Use Case:*
 - ▶ Caller is within parallel kernel body with a batch of input data
 - ▶ *Multiple Interfaces: Serial, Team, TeamVector*
 - ▶ Serial: no nested parallelism is used internally
 - ▶ Team: one-level nested parallelism is used with *TeamThreadRange*
 - ▶ TeamVector: two-level nested parallelism is used with *TeamThreadRange* and *TeamVectorRange*

Batched BLAS/LAPACK is **simple** i.e., BLAS/LAPACK in a parallel loop

```
auto A = Kokkos::View<double***>('A', N, Blk, Blk);  
Kokkos::parallel_for( RangePolicy(N), /// users' parallel execution policy  
  KOKKOS_LAMBDA(int &i) {  
    auto AA = Kokkos::subview(A, i, ALL, ALL);  
    KokkosBatched::SerialLU(AA); /// functor-level interface  
  });
```

Kokkos batched BLAS/LAPACK is made up of following two components

- ▶ Kokkos parallel execution policy with `parallel_for`
- ▶ A functor-level interface to be used in `operator()`

Hierarchical functor interface is required to match Kokkos' hierarchical parallelism

Device Interface

- ▶ internally uses TeamPolicy
- ▶ is used for large input data that occupies an entire device
- ▶ can use an execution space instance to launch in a stream

Device with ExecutionSpace

```
Kokkos::Cuda execution_space(myCudaStream);  
KokkosBlas(execution_space);
```


TeamVector Interface

- ▶ internally uses two nested `parallel_for` with `TeamThreadRange` and `ThreadVectorRange`
- ▶ requires the member (thread communicator) as an input argument

TeamVector with TeamPolicy

```
parallel_for(TeamPolicy,
  KOKKOS_LAMBDA(member_type &member){
    KokkosBatched::TeamVectorDoSomething(member);
  });
```

Team Interface

- ▶ internally uses TeamThreadRange only
- ▶ in general is used with SIMD or Ensemble types where vector parallelism is expressed within the type
- ▶ can include ThreadVectorRange

Team without ThreadVectorRange

```
parallel_for( TeamPolicy ,  
  KOKKOS_LAMBDA( member_type &member){  
    KokkosBatched :: TeamDoThing( member);  
  });
```

Team with ThreadVectorRange outside

```
parallel_for( TeamPolicy ,  
  KOKKOS_LAMBDA( member_type &member){  
    parallel_for( ThreadVectorRange) {  
      KokkosBatched :: TeamDoSomething(  
        member);  
    }); };
```

Serial Interface

- ▶ can be used in a flat `parallel_for` i.e., `Kokkos::RangePolicy`
- ▶ can be used in the most inner loop of nested `parallel_for`'s

Serial with RangePolicy

```
parallel_for(RangePolicy ,  
  KOKKOS_LAMBDA(int &idx){  
    KokkosBatched::SerialDoThing();  
  });
```

Serial in Hierarchical parallel loops

```
parallel_for(TeamPolicy ,  
  KOKKOS_LAMBDA(member_type &member){  
    parallel_for(TeamThreadRange) {  
      parallel_for(ThreadVectorRange) {  
        KokkosBatched::  
          SerialDoSomething();  
      };  
    };  
  });
```

Summary: Batched BLAS/LAPACK

- ▶ User composable (batched) BLAS interface: parallel execution policy + functor-level interface
- ▶ Performance on GPUs is tunable:
 - ▶ Launching light-weight kernels multiple times can cause overhead
 - ▶ Fusing too many functor-level BLAS/LAPACK operations is difficult to do while maintaining optimal performance with a single team size