

# The Kokkos Lectures

## Kokkos Execution Space Instances

May 5, 2022

Sandia National Laboratories is a multi-mission laboratory managed and operated by National Technology and Engineering Solutions of Sandia, LLC., a wholly owned subsidiary of Honeywell International, Inc., for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA0003525.  
SAND2020-8508 PE

## **This is NOT a general tutorial!**

Today: three topics with exercises.

- ▶ Three Short Lectures
  - ▶ Execution Space Instances
  - ▶ PyKokkos
  - ▶ New KokkosKernels Capabilities
- ▶ Hands-On Exercise
  - ▶ Split Into Breakout Rooms
  - ▶ One per topic: Choose which one you want to attend

Tuesday 2022-05-10: Community BOF Days with presentations on more new capabilities.

<https://www.exascaleproject.org/event/ecp-community-bof-days-2022>

## Online Resources:

- ▶ <https://github.com/kokkos>:
  - ▶ Primary Kokkos GitHub Organization
- ▶ <https://github.com/kokkos/kokkos-tutorials/wiki/Kokkos-Lecture-Series>:
  - ▶ Slides, recording and Q&A for the Lectures
- ▶ <https://github.com/kokkos/kokkos/wiki>:
  - ▶ Wiki including API reference
- ▶ <https://kokkosteam.slack.com>:
  - ▶ Slack channel for Kokkos.
  - ▶ Please join: fastest way to get your questions answered.
  - ▶ Can whitelist domains, or invite individual people.

Instructions for joining the exercises can be found here:

<https://github.com/kokkos/kokkos-tutorials/issues/60>

Issue: ECP Annualmeeting 2022 - Kokkos Tutorial AWS Instances

Claim an instance in the Google spreadsheet.

SSH to the instance.

The slides are available at:

<https://github.com/kokkos/kokkos-tutorials>

under Other/ECP-Annualmeeting/2022

# Asynchronicity and ExecutionSpace Instances

The (non-)blocking behavior of Kokkos operations.

## Learning objectives:

- ▶ What are blocking and non-blocking operations in Kokkos.
- ▶ What kind of work can overlap.
- ▶ How to wait for completion.
- ▶ How to run kernels simultaneously on a GPU.

### **Most operations in Kokkos are non-blocking**

- ▶ The caller returns before the operation is finished
- ▶ The caller can do other things, while operations are executing

### **Most operations in Kokkos are non-blocking**

- ▶ The caller returns before the operation is finished
- ▶ The caller can do other things, while operations are executing

### **So what is the ordering behavior?**

- ▶ Execution Spaces have an ordered execution queue
- ▶ The queue is first-in/first-out (FIFO)

### **Most operations in Kokkos are non-blocking**

- ▶ The caller returns before the operation is finished
- ▶ The caller can do other things, while operations are executing

### **So what is the ordering behavior?**

- ▶ Execution Spaces have an ordered execution queue
- ▶ The queue is first-in/first-out (FIFO)

### **Important Point**

Execution Spaces execute operations in dispatch order.



## Execution Space Instances

- ▶ Each unique **Instance** of an execution space has its own queue
- ▶ Execution Policies can take an instance as the first argument
- ▶ `deep_copy` can take an instance as a first argument
- ▶ For every **Execution Space Type** there is a **default instance**
  - ▶ It is a singleton
  - ▶ The **default instance** is returned by the default constructor
  - ▶ Used if no specific instance is provided

## Execution Space Instances

- ▶ Each unique **Instance** of an execution space has its own queue
- ▶ Execution Policies can take an instance as the first argument
- ▶ `deep_copy` can take an instance as a first argument
- ▶ For every **Execution Space Type** there is a **default instance**
  - ▶ It is a singleton
  - ▶ The **default instance** is returned by the default constructor
  - ▶ Used if no specific instance is provided

```
// This is equivalent:  
RangePolicy<ExecSpace>  
  policy_1(0, N);  
RangePolicy<ExecSpace>  
  policy_2(ExecSpace(), 0, N);
```

## We use the following conventions in subsequent slides

```
// Execution Space types
```

```
using device = Kokkos::DefaultExecutionSpace;  
using host = Kokkos::DefaultHostExecutionSpace;
```

```
// Execution Space instances
```

```
device dev1(..), dev2(..)  
device host1(..), host2(..);
```

```
// Execution Policies
```

```
RangePolicy<device> policy_d(0,N), policy_device(0,N);  
RangePolicy<host> policy_h(0,K), policy_host(0,K);  
RangePolicy<device> policy_d1(dev1, 0,N), policy_d2(dev2, 0,N);  
RangePolicy<host> policy_h1(host1, 0,N), policy_h2(host2, 0,N);
```

```
// Functors/Lambda for parallel_for
```

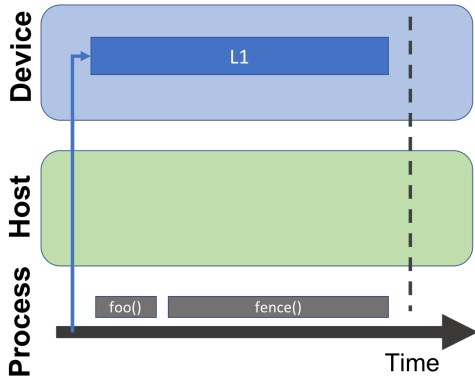
```
auto L1 = KOKKOS_LAMBDA(int i) {...};  
auto L2 = ...; auto L3 = ...; auto L4 = ...; auto L5 = ...;
```

```
// Functors/Lambda for parallel_reduce
```

```
auto R1 = KOKKOS_LAMBDA(int i, double& lsum) {...};
```

## Most Kokkos Operations are Asynchronous


- ▶ Best to assume all of them are asynchronous
- ▶ They overlap with work in the process thread
- ▶ Use `Kokkos::fence()` to wait for completion



```
RangePolicy<>  
    policy_device(0,N)  
FunctorL1 L1(...);  
  
parallel_for("L1",  
    policy_device, L1);  
foo();  
fence();
```

## Execution Spaces execute in dispatch order

- ▶ Dispatches to the same space instance will never overlap
- ▶ Executed in order FIFO
- ▶ Use `Kokkos::fence()` to wait for completion




```
RangePolicy<>
    policy_device(0,N)
FunctorL1 L1(...);
FunctorL2 L2(...);

parallel_for("L1",
    policy_device, L1);
parallel_for("L2",
    policy_device, L2);
foo();
fence();
```

## Execution Spaces execute in dispatch order

- ▶ Dispatches to the same space instance will never overlap
- ▶ Executed in order FIFO
- ▶ Use `Kokkos::fence()` to wait for completion




```
RangePolicy<>
    policy_device(0,N)
FunctorL1 L1(...);
FunctorL2 L2(...);

parallel_for("L1",
    policy_device, L1);
foo();
parallel_for("L2",
    policy_device, L2);
fence();
```

## ExecutionSpaces are Independent

- ▶ Dispatches into different ExecutionSpaces may overlap.
- ▶ Overlap with process thread functions and each other
- ▶ Use `Kokkos::fence()` to wait for completion of all




```

RangePolicy<>
    policy_d(0,N)
RangePolicy<Host>
    policy_host(0,N)
FunctorL1 L1(...);
FunctorL2 L2(...);
FunctorL3 L3(...);
parallel_for("L1",
    policy_d, L1);
parallel_for("L2",
    policy_host, L2);
parallel_for("L3",
    policy_d, L3);
foo();
fence();
  
```

## Reality: Some Host Backends Block

- ▶ Most host backends are blocking dispatches (except HPX)
- ▶ They never overlap with process thread functions
- ▶ But: **Do NOT** rely on blocking behavior!!




```
RangePolicy<>
    policy_d(0,N)
RangePolicy<Host>
    policy_host(0,N)
FunctorL1 L1(...);
FunctorL2 L2(...);
FunctorL3 L3(...);
parallel_for("L1",
    policy_d, L1);
parallel_for("L2",
    policy_host, L2);
parallel_for("L3",
    policy_d, L3);
foo();
fence();
```



## Reductions to Scalars are Blocking

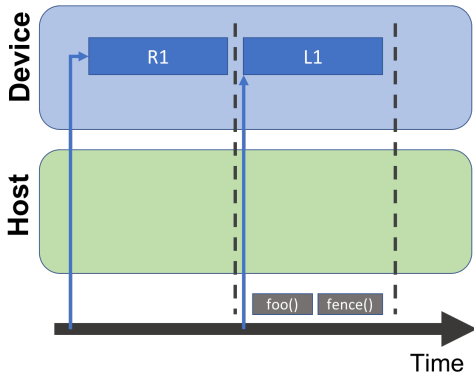
- ▶ The call only returns after result is available.
- ▶ FIFO implies, every other kernel in the same ExecutionSpace instance will be done too



```
RangePolicy<>  
    policy_d(0,N)  
FunctorL1 L1(...);  
FunctorR1 R1(...);  
  
double result;  
parallel_for("L1",  
    policy_d, L1);  
parallel_reduce("R1",  
    policy_d, R1, result);  
foo();
```

## Reductions to Scalars are Blocking

- ▶ The call only returns after result is available.
- ▶ For subsequent dispatches previous rules apply

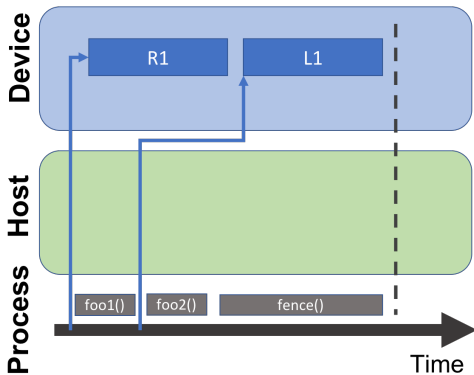


```
RangePolicy<>
    policy_d(0,N)
FunctorL1 L1(...);
FunctorR1 R1(...);

double result;
parallel_reduce("R1",
    policy_d, R1, result);
parallel_for("L1",
    policy_d, L1);
foo();
fence();
```

## Reductions to Views are Non-blocking


- ▶ Behave like a `parallel_for`
- ▶ Results are only available after a `Kokkos::fence()`
- ▶ Even true for unmanaged Views of host variables!



```
...
double result;
View<double, HostSpace>
    v_result(&result);
parallel_reduce("R1",
    policy_d, R1, v_result);
foo1();
parallel_for("L1",
    policy_d, L1);
foo2();
fence();
```

## Simple Parallel Loop


- ▶ Asynchronous
- ▶ Overlaps with host functions
- ▶ Use `Kokkos::fence()` to wait for completion



```
RangePolicy<>  
    policy_device(0,N)  
FunctorL1 L1(...);  
FunctorL2 L2(...);  
  
parallel_for("L1",  
    policy_device, L1);  
parallel_for("L2",  
    policy_device, L2);  
foo();  
fence();
```

## 2-Argument `deep_copy` is fully blocking

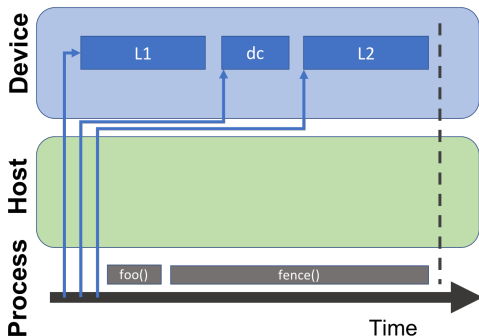
- ▶ Implies a full fence before the copy
- ▶ Copy is done by the time call returns.
- ▶ Even if it is a no-op due to `src == dst`!



```
parallel_for("L1",  
            policy_device, L1);  
deep_copy(dest,src);  
parallel_for("L2",  
            policy_device, L2);  
foo();  
fence();
```

## deep\_copy with space argument are non-blocking

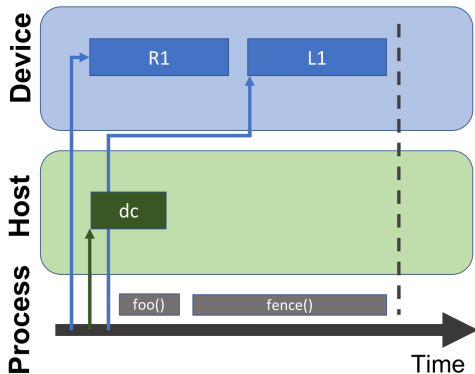
- ▶ Execute in dispatch order in the queue of the space
- ▶ Overlap with host process functions
- ▶ Use `Kokkos::fence()` to wait for completion



```
parallel_for("L1",  
            policy_device, L1);  
deep_copy(device,  
          dest,src);  
parallel_for("L2",  
            policy_device, L2);  
foo();  
fence();
```

## deep\_copy with space argument are non-blocking

- ▶ Execute in dispatch order in the queue of the space
- ▶ Overlap with other execution spaces
- ▶ Use `Kokkos::fence()` to wait for completion



```
parallel_for("L1",  
            policy_device, L1);  
deep_copy(host(),  
          dest,src);  
parallel_for("L2",  
            policy_device, L2);  
foo();  
fence();
```

## So what about **CUDA/HIP streams** and **SYCL queues**?

Up to now we only used default execution space instances, but what if you want to have concurrent kernels on the GPU?



### **So what about CUDA/HIP streams and SYCL queues?**

Up to now we only used default execution space instances, but what if you want to have concurrent kernels on the GPU?

#### Execution Space Instances

Execution Space instances behave largely like CUDA/HIP streams

## So what about CUDA/HIP streams and SYCL queues?

Up to now we only used default execution space instances, but what if you want to have concurrent kernels on the GPU?

### Execution Space Instances

Execution Space instances behave largely like CUDA/HIP streams

You can create different instances:

```
ExecSpace space;  
std::vector<ExecSpace> instances =  
    Kokkos::Experimental::partition_space(space, w1, ..., wN);
```

- Creates a vector of N new instances of ExecSpace

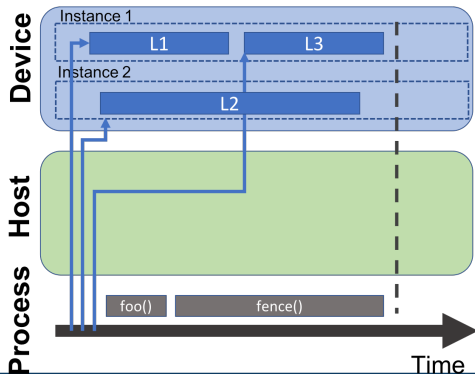
## So what about CUDA/HIP streams and SYCL queues?

```
ExecSpace space;  
std::vector<ExecSpace> instances =  
    Kokkos::Experimental::partition_space(space, w1, ..., wN);
```

- ▶ The new instances will submit to a subset of the resources associated with space.
- ▶ The weights are hints as to the fraction of resources these instances are allowed to dispatch too.
  - ▶ Use integers or floating point numbers
  - ▶ Number of instances created is equal to number of weights
  - ▶ Note weights are hints, and might be ignored.
- ▶ Returns a `std::vector` of execution space instances
- ▶ The function may return just N copies of space
- ▶ all the created instances are independent of each other and from the original instance

## Instances of Execution Spaces own an exec queue

- ▶ Work dispatched to different instances overlaps with each other
- ▶ Overlaps with host process functions
- ▶ Use `Kokkos::fence()` to wait for completion of all




```
auto instances = partition
Device dev1 = instances[0];
Device dev2 = instances[1];
RangePolicy<Device>
    policy_d1(dev1,0,N);
RangePolicy<Device>
    policy_d2(dev2,0,N);
```

```
parallel_for("L1",
    policy_d1, L1);
parallel_for("L2",
    policy_d2, L2);
parallel_for("L3",
    policy_d1, L3);
foo();
fence();
```

## deep\_copy with an instance argument also overlap

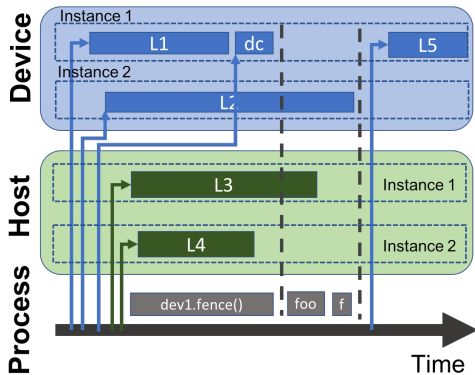
- ▶ deep\_copy with an instance argument are like any other parallel operation
- ▶ Overlaps with parallel operations in other instance



```
parallel_for("L1",  
    policy_d1, L1);  
parallel_for("L2",  
    policy_d2, L2);  
deep_copy(dev1,  
    dest, src);  
parallel_for("L3",  
    policy_d1, L3);  
foo();  
fence();
```

## There are instance fences


- ▶ Use instance specific fence to only wait on that instance
- ▶ Operations in other instances can overlap with that fence
- ▶ Use `Kokkos::fence()` to wait for all outstanding ops



```
parallel_for("L1",
    policy_d1, L1);
parallel_for("L2",
    policy_d2, L2);
deep_copy(dev1,
    dest, src);
parallel_for("L3",
    policy_h1, L3);
parallel_for("L4",
    policy_h2, L4);
dev1.fence();
foo();
fence();
parallel_for("L5",
    policy_d1, L5);
```

## Reality Check: Kokkos Views deallocation implies fence!

- ▶ Due to limitations of reference counting, deallocations fence!
- ▶ **Important:** this is implementation limitation not semantic!
- ▶ **Do NOT** rely on deallocations fencing!



```

{
  View<...> tmp(...);
  parallel_for("L1",
    policy_d1, L1);
  parallel_for("L2",
    policy_d2, L2);
  parallel_for("L3",
    policy_h1, L3);
  parallel_for("L4",
    policy_h2, L4);
}
foo();
fence();
parallel_for("L5",
  policy_d1, L5);

```

- ▶ Execution Space Instances execute work in order of dispatch.
- ▶ Operations dispatched to different Execution Space Instances can overlap.
- ▶ Each Execution Space type has a default instance as a singleton.
- ▶ Use `Kokkos::fence()` to wait for completion of ALL outstanding work.
- ▶ Use `exec_space_instance.fence()` to wait for completion of outstanding work dispatched to a specific execution space instance.



**Exercise:** Explore impact of using instances for concurrent execution.

Getting set up in your home directory:

```
mkdir Kokkos  
cd Kokkos  
git clone https://github.com/kokkos/kokkos  
git clone https://github.com/kokkos/kokkos-tutorials
```

Find the exercise in the kokkos-tutorials/Exercises/instances folder.

The Begin subdirectory contains the code. Only instances\_begin.cpp needs modifications.

Simply type `make -j 8` to build the exercise.

Look for EXERCISE comments to find places to modify.

**Tasks:**

- ▶ Create ExecutionSpace instances
- ▶ Use the instances to dispatch the kernels and `deep_copy`
- ▶ Remember to use device and host accessible views as results for reductions!

**Things to try:**

- ▶ Compare using the same instance for both data sets, and different ones.
- ▶ Compare performance for small (e.g. -N 5000) and large (e.g. -N 30000) problems.
- ▶ What happens if you create the result Views inside the repeat loop?